


Product variability — two real shapes, as instances (not classes)




Product-level-only image.
Normal, not orphaned — a
variation never has to
claim every product image.

Computed, not stored — this is the other thing a class diagram cannot show. `price_override=null` means `ResolveVariationPrice::current()` returns the `PRODUCT`'s `regular_price` (20.00) at read time. Nothing here is copied or kept in sync; there is no second value to go stale.

A third, unrelated table. Free-text name/value on the PRODUCT, not a pivot, not per-variation, not filterable — "Cotton", "cotton" and "100% Cotton" are three ungroupable values no query can filter on. The right home for prose a customer reads but nobody searches by. See [explanation/product-variability.md](#). "Why not product_specifications for the descriptive half."

Plain Mug — no variations beyond one, NO images anywhere



The one distinction worth holding onto: Red and S are BOTH AttributeValue rows — one vocabulary table, used by two DIFFERENT pivots for two DIFFERENT questions. A class diagram would show one association and hide this; the green vs. blue edges here are the same table used twice, on purpose.